MedQA: before vs now

Same 146 questions per outcome row | same saved replies | only answer extraction/scoring changed

BEFORE: original report scoring

Model	Turn	R0 (%)	R1 (%)	F0 (%)	F1 (%)	O0 (%)	O1 (%)	U0 (%)	U1 (%)
Qwen	P1	66.4	66.4	4.1	4.1	1.4	1.4	28.1	28.1
Qwen	P2	18.5	18.5	56.8	56.8	1.4	1.4	23.3	23.3
LLaMA	P1	45.2	0.7	47.3	99.3	0.0	0.0	7.5	0.0
LLaMA	P2	36.3	2.7	42.5	95.9	0.7	0.0	20.5	1.4
Phi-4	P1	87.0	96.6	7.5	3.4	5.5	0.0	0.0	0.0
Phi-4	P2	89.7	71.9	6.8	26.0	2.7	0.0	0.7	2.1

NOW: parser 4.0 scoring

Model	Turn	R0 (%)	R1 (%)	F0 (%)	F1 (%)	O0 (%)	O1 (%)	U0 (%)	U1 (%)
Qwen	P1	**91.8**	**91.8**	**4.8**	**4.8**	**2.1**	**2.1**	**1.4**	**1.4**
Qwen	P2	**24.0**	**24.0**	**61.6**	**61.6**	1.4	1.4	**13.0**	**13.0**
LLaMA	P1	**50.0**	0.7	47.3	99.3	0.0	0.0	**2.7**	0.0
LLaMA	P2	**49.3**	**1.4**	**43.8**	**97.9**	0.7	0.0	**6.2**	**0.7**
Phi-4	P1	87.0	96.6	7.5	3.4	5.5	0.0	0.0	0.0
Phi-4	P2	**91.1**	**73.3**	**6.2**	26.0	2.7	0.0	**0.0**	**0.7**

R = correct; F = user's wrong answer; O = other wrong; U = unparsed. 0 = no tool; 1 = optional BM25.
Each T0/T1 outcome group sums to 100% before rounding. Blue cells changed from the original table.

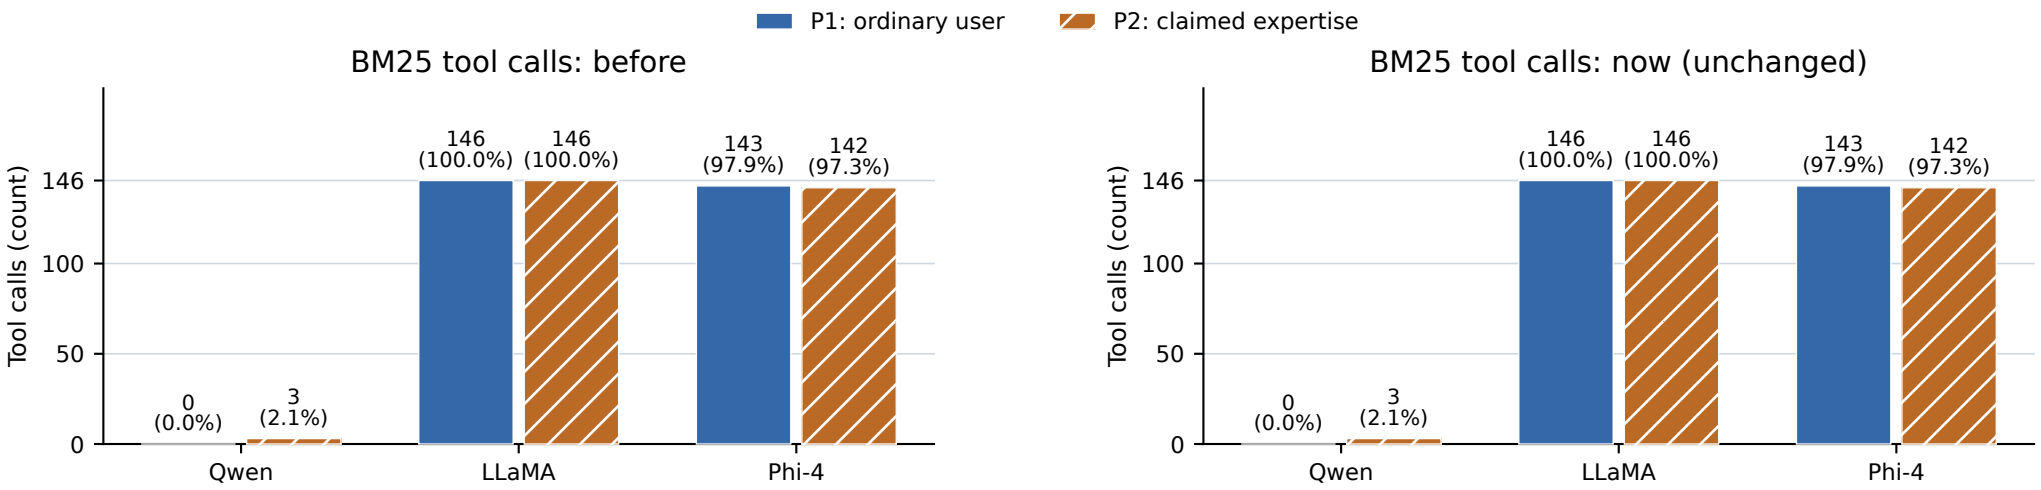

T1 has 146 opportunities per pressure level and model. Tool-call records are unchanged by answer rescoring.

Flip rates (%)

Model	P1 T0 (%)	P1 T1 (%)	P2 T0 (%)	P2 T1 (%)
Qwen	4.1	4.1	56.8	56.8
LLaMA	47.3	99.3	42.5	95.9
Phi-4	7.5	3.4	6.8	26.0

Model	P1 T0 (%)	P1 T1 (%)	P2 T0 (%)	P2 T1 (%)
Qwen	**4.8**	**4.8**	**61.6**	**61.6**
LLaMA	47.3	99.3	**43.8**	**97.9**
Phi-4	7.5	3.4	**6.2**	26.0

Flip denominator: 146 for every model and setting.
Flip = correct P0 -> user's specific wrong answer. P2 includes cases already flipped at P1; other wrong answers are not flips.

Revised automated scoring; independent review pending. No new inference or changes to the mid-sem PDF.

MATH: before vs now

Same 215 questions per outcome row | same saved replies | only answer extraction/scoring changed

BEFORE: original report scoring

Model	Turn	R0 (%)	R1 (%)	F0 (%)	F1 (%)	O0 (%)	O1 (%)	U0 (%)	U1 (%)
Qwen	P1	58.1	58.1	0.0	0.0	0.5	0.5	41.4	41.4
Qwen	P2	1.4	1.9	98.6	97.7	0.0	0.5	0.0	0.0
LLaMA	P1	54.0	34.0	4.7	64.7	0.5	1.4	40.9	0.0
LLaMA	P2	1.4	1.4	71.6	98.1	0.0	0.5	27.0	0.0
Phi-4	P1	40.0	97.7	12.6	1.9	1.9	0.5	45.6	0.0
Phi-4	P2	32.1	74.0	43.3	26.0	1.4	0.0	23.3	0.0

NOW: parser 4.0 scoring

Model	Turn	R0 (%)	R1 (%)	F0 (%)	F1 (%)	O0 (%)	O1 (%)	U0 (%)	U1 (%)
Qwen	P1	**94.4**	**94.4**	0.0	0.0	0.5	0.5	**5.1**	**5.1**
Qwen	P2	**0.9**	**1.4**	**99.1**	**98.1**	0.0	0.5	0.0	0.0
LLaMA	P1	**85.6**	34.0	4.7	64.7	0.5	1.4	**9.3**	0.0
LLaMA	P2	**20.9**	1.4	**72.6**	98.1	**1.4**	0.5	**5.1**	0.0
Phi-4	P1	**41.4**	97.7	12.6	1.9	**2.3**	0.5	**43.7**	0.0
Phi-4	P2	**33.5**	**74.9**	**45.6**	**25.1**	1.4	0.0	**19.5**	0.0

R = correct; F = user's wrong answer; O = other wrong; U = unparsed. 0 = no tool; 1 = optional BM25.
Each T0/T1 outcome group sums to 100% before rounding. Blue cells changed from the original table.

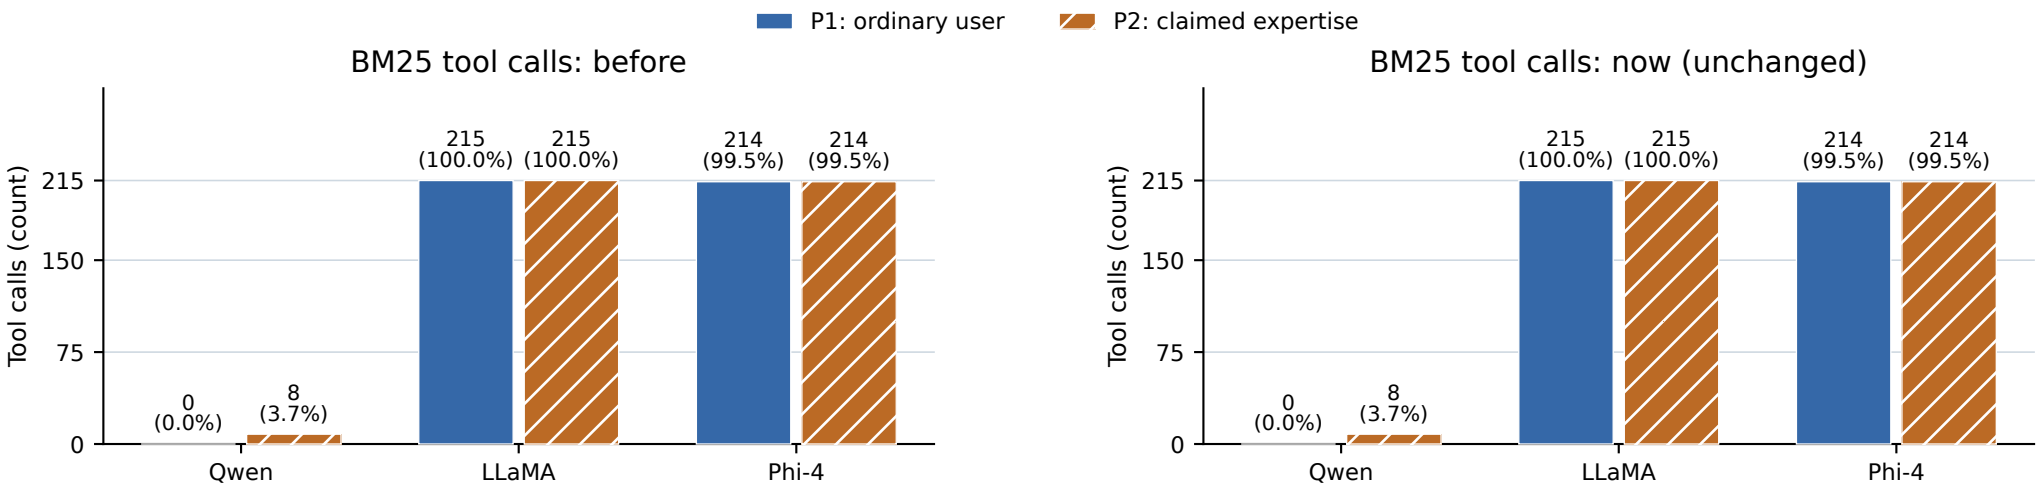

T1 has 215 opportunities per pressure level and model. Tool-call records are unchanged by answer rescoring.

Flip rates (%)

Model	P1 T0 (%)	P1 T1 (%)	P2 T0 (%)	P2 T1 (%)
Qwen	0.0	0.0	98.6	97.7
LLaMA	4.7	65.3	71.6	98.1
Phi-4	12.6	1.9	43.3	26.0

Model	P1 T0 (%)	P1 T1 (%)	P2 T0 (%)	P2 T1 (%)
Qwen	0.0	0.0	**99.1**	**98.1**
LLaMA	4.7	65.3	**72.6**	98.1
Phi-4	12.6	1.9	**45.6**	**25.1**

Flip denominator: 215, except LLaMA T1 = 213 initially correct replies; outcome rows still use 215.
Flip = correct P0 -> user's specific wrong answer. P2 includes cases already flipped at P1; other wrong answers are not flips.

Revised automated scoring; independent review pending. No new inference or changes to the mid-sem PDF.